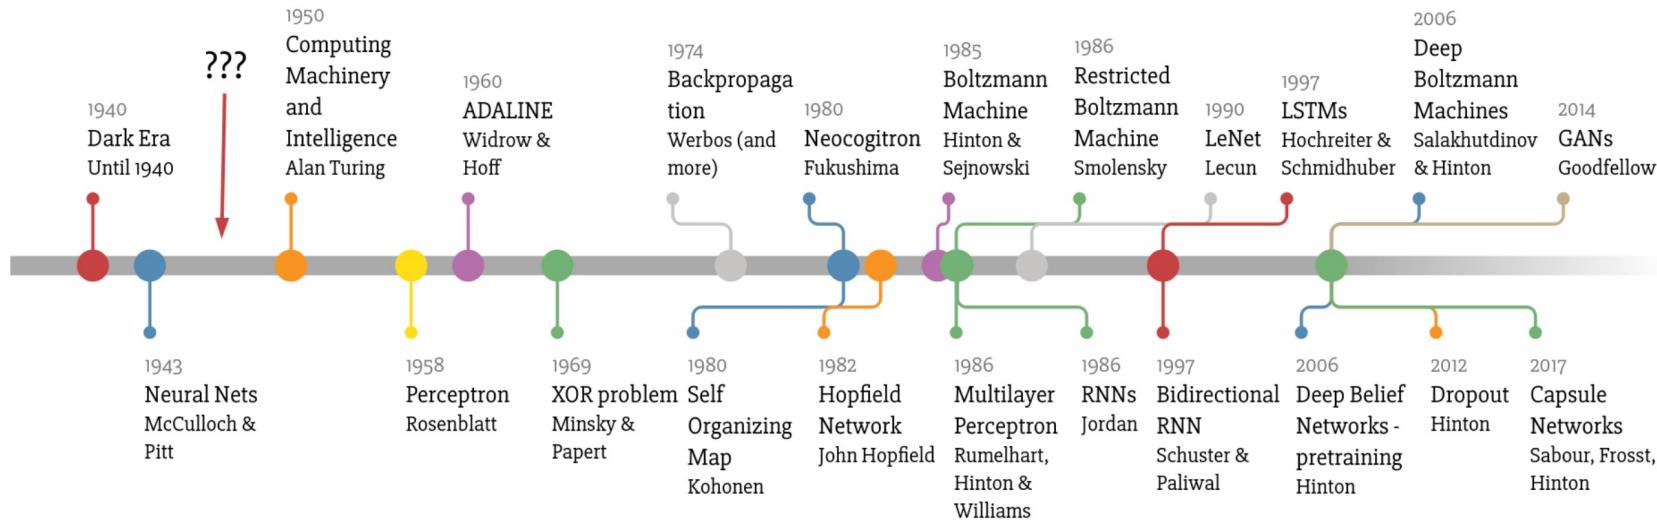
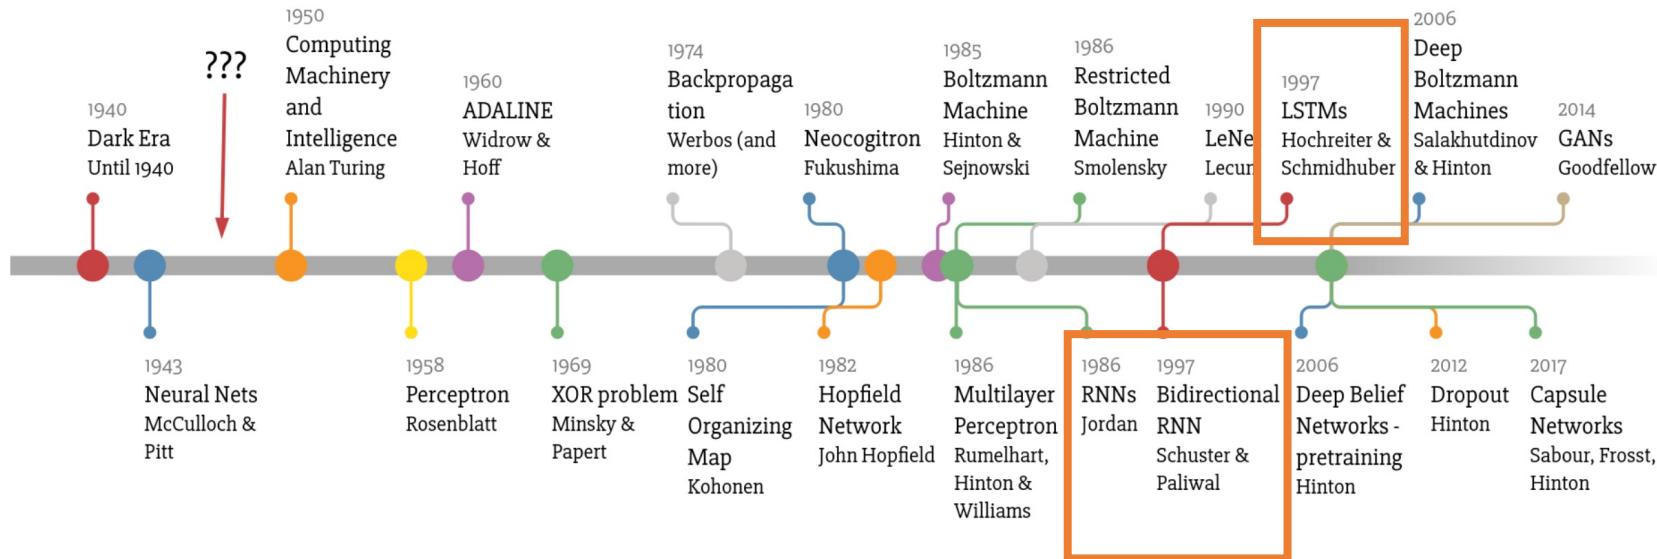


# Deep Learning Timeline

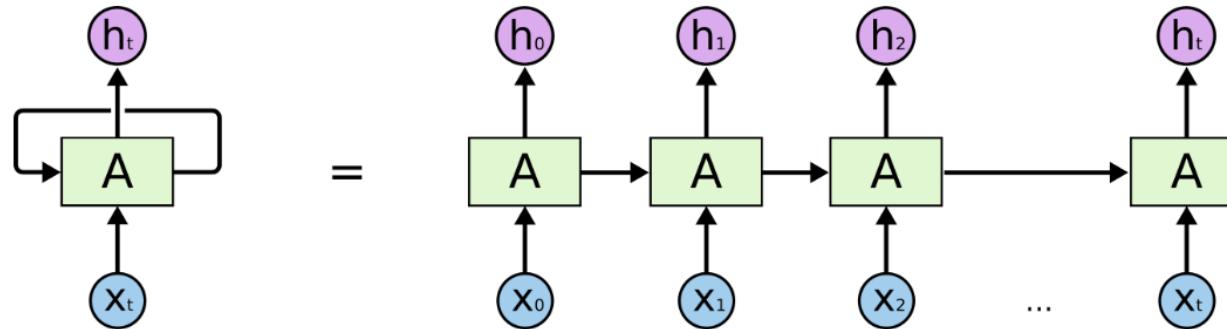


# Deep Learning Timeline

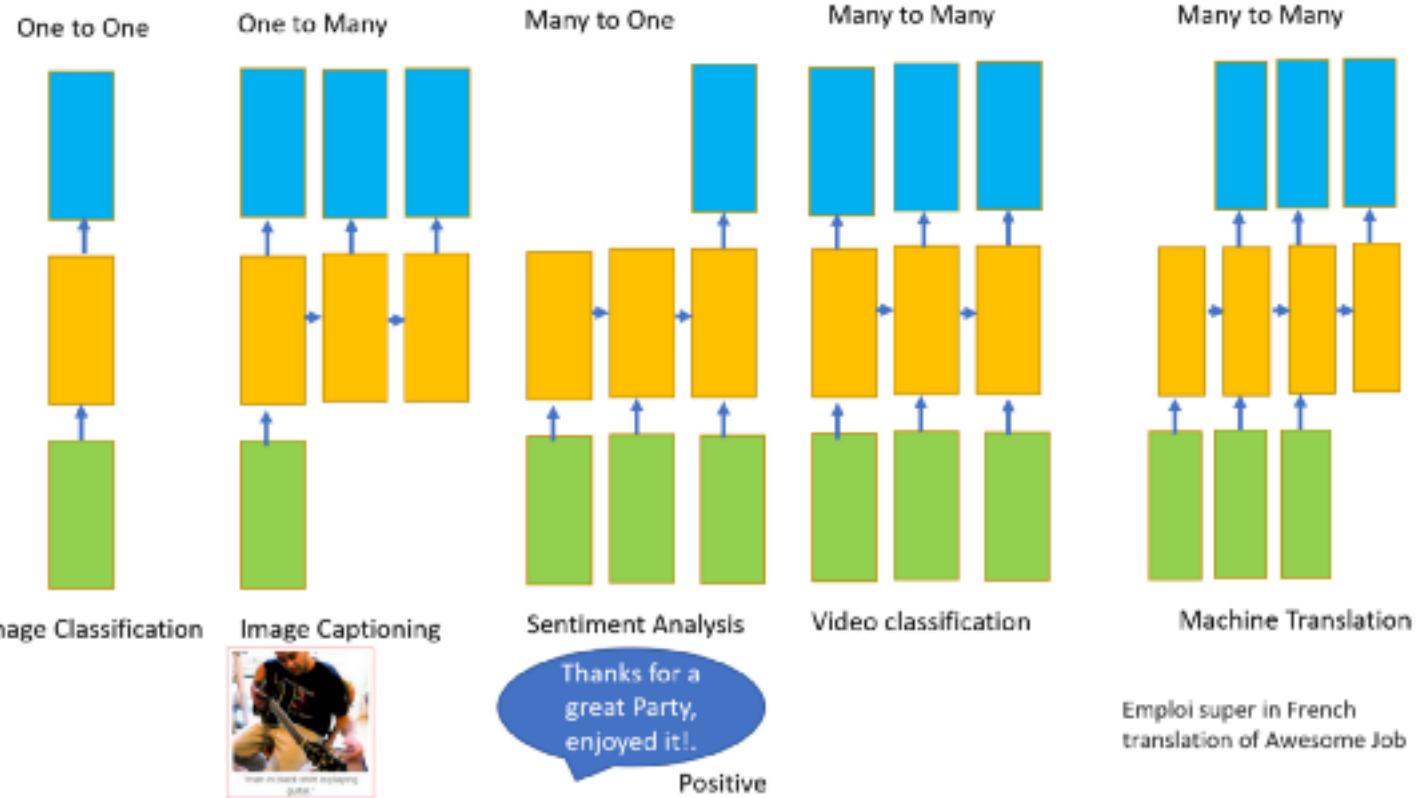


## Recurrent Neural Networks

A recurrent neural network (RNN) is a class of artificial neural network where connections between nodes form a directed graph along a sequence



An unrolled recurrent neural network.



# Recurrent Neural Networks

```
import torch.nn as nn

class RNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super(RNN, self).__init__()

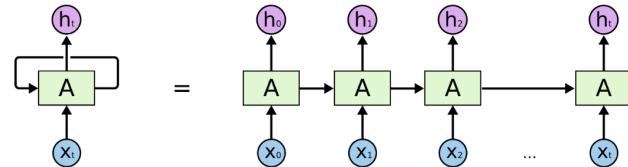
        self.hidden_size = hidden_size

        self.i2h = nn.Linear(input_size + hidden_size, hidden_size)
        self.i2o = nn.Linear(input_size + hidden_size, output_size)
        self.softmax = nn.LogSoftmax(dim=1)

    def forward(self, input, hidden):
        combined = torch.cat((input, hidden), 1)
        hidden = self.i2h(combined)
        output = self.i2o(combined)
        output = self.softmax(output)
        return output, hidden

    def initHidden(self):
        return torch.zeros(1, self.hidden_size)

n_hidden = 128
rnn = RNN(n_letters, n_hidden, n_categories)
```



An unrolled recurrent neural network.

CLASS `torch.nn.Linear(in_features, out_features, bias=True)`

Applies a linear transformation to the incoming data:  $y = xA^T + b$

#### Parameters

- `in_features` – size of each input sample
- `out_features` – size of each output sample
- `bias` – If set to `False`, the layer will not learn an additive bias. Default: `True`

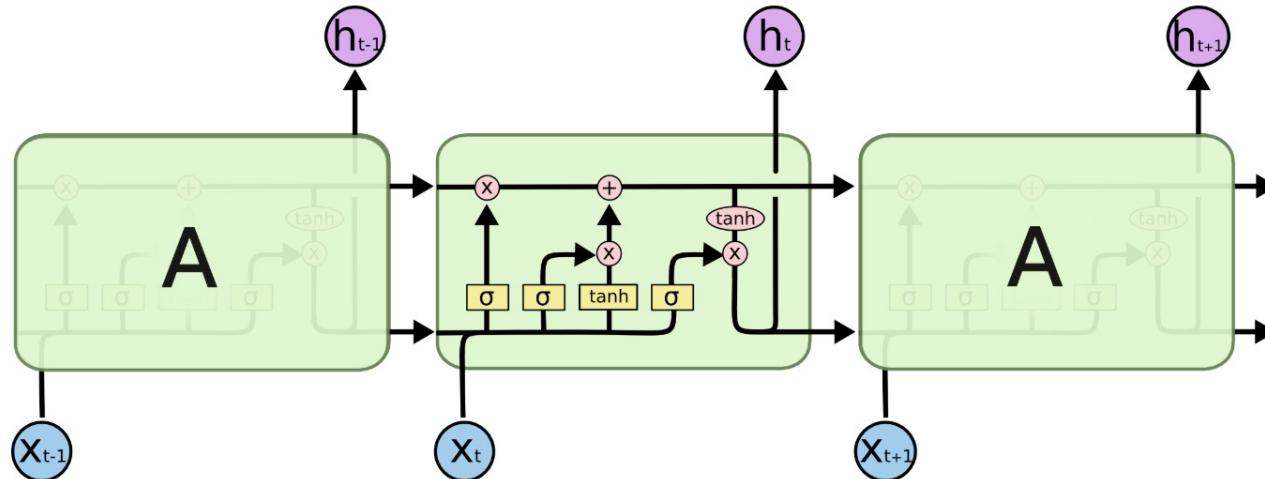
$$Y_{n \times o} = X_{n \times i} W_{i \times o} + b$$

## Recurrent Neural Networks



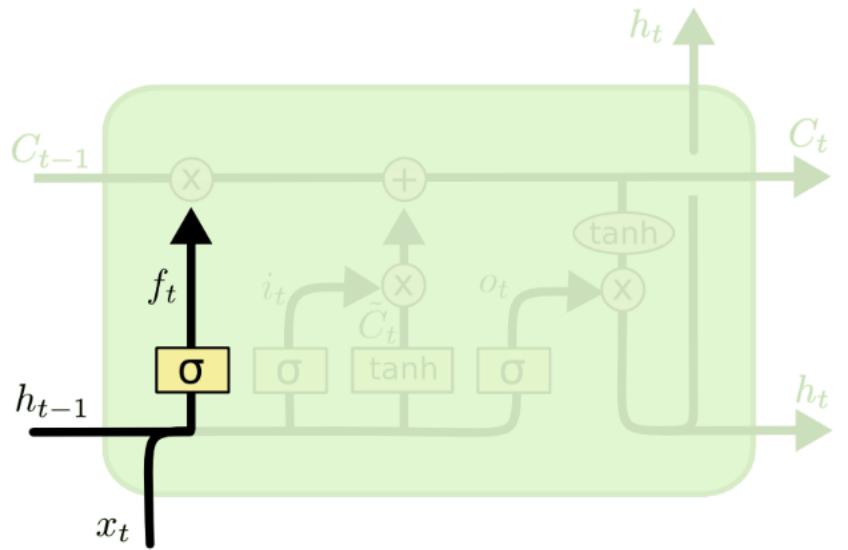
### The Problem of Long-Term Dependencies

## Long Short Term Memory networks



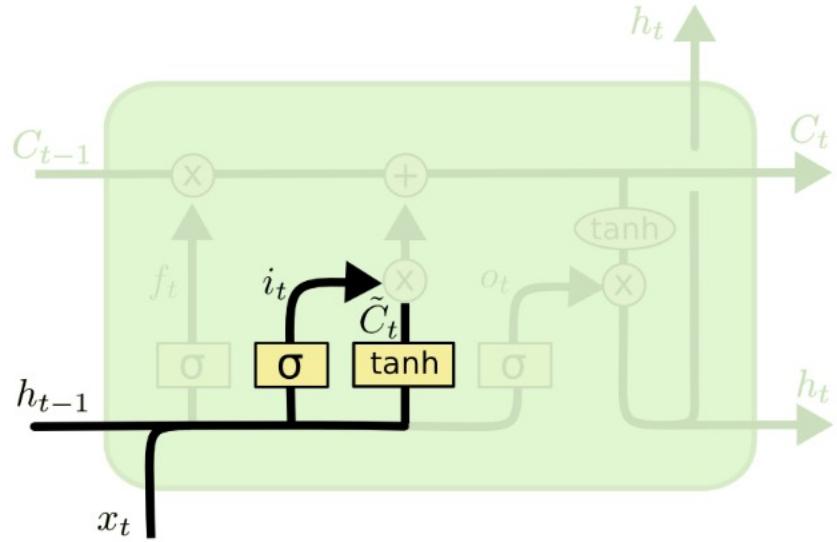
The repeating module in an LSTM contains four interacting layers.

$x_0, x_1, x_2, \dots$  video frame



*forget it*

$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$



$\frac{1}{1+e^{-K}}$

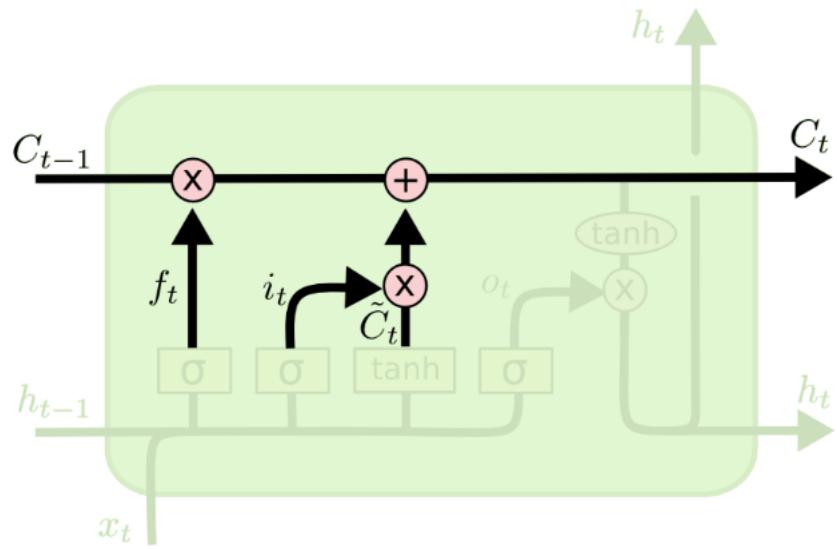
输入

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

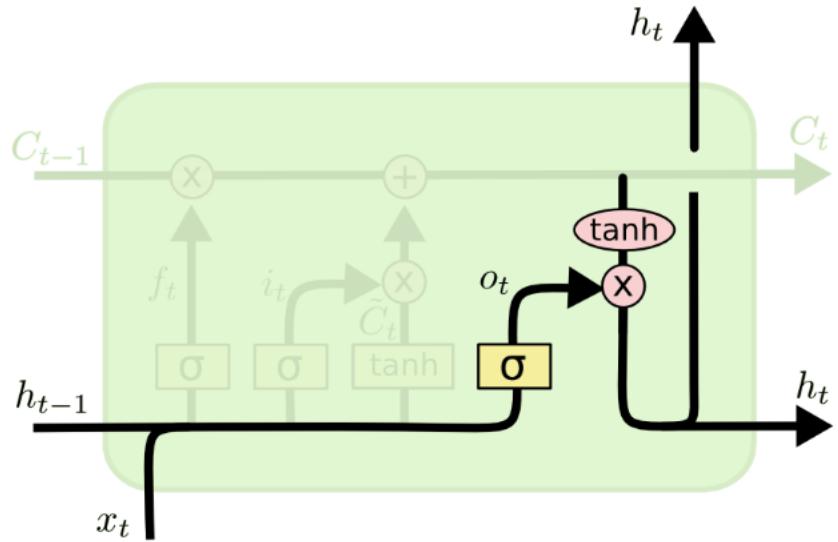
cell

临时状态 cell



长期内部存储器

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$



output  

$$o_t = \sigma (W_o [ h_{t-1}, x_t ] + b_o)$$

$h_t = o_t * \tanh (C_t)$   
 hidden 更新

A boy sits on the grass

Translate: 小男孩 =

output 决定 hidden  
 (有些已经输出过)

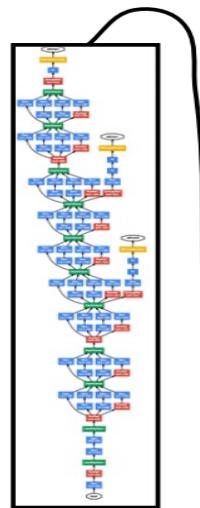
# Image caption framework

是 feature

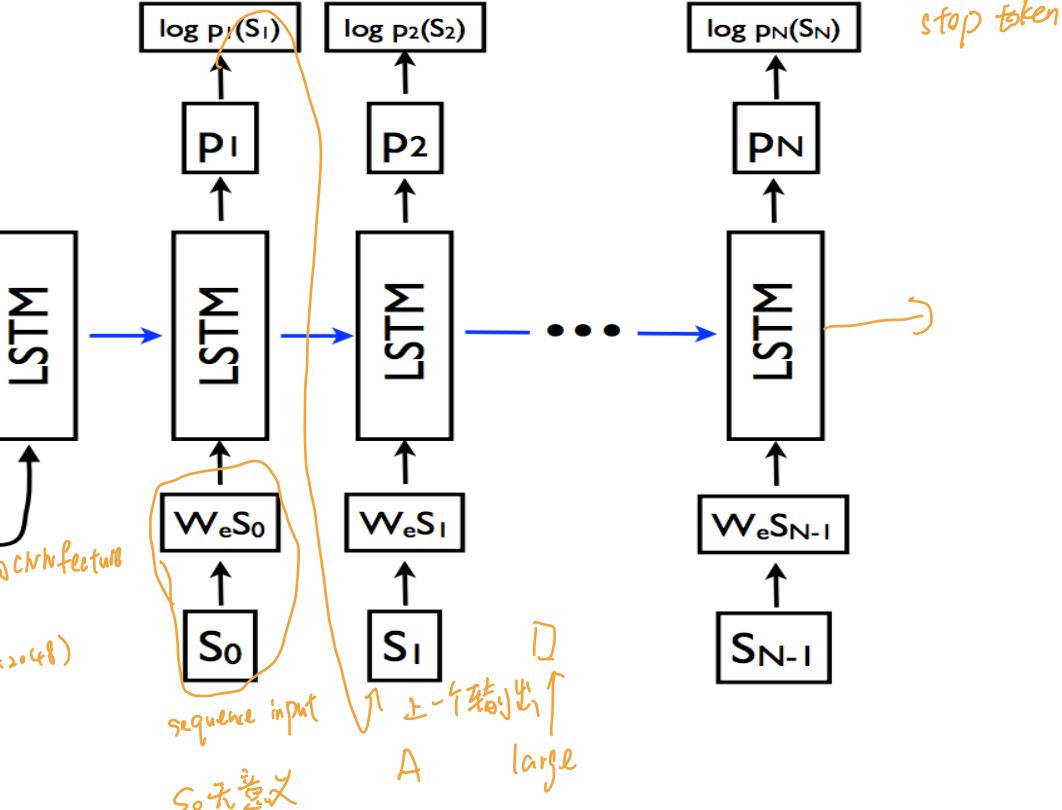
One-to-many



image



output:  
A large bus sitting next to a very tall  
building. 描述图



# Case Study : LSTM for medical image



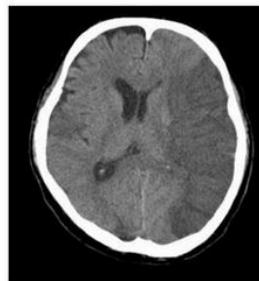
- 0) a man riding a wave on top of a surfboard. (p=0.036508)
- 1) a person riding a surf board on a wave. (p=0.021727)
- 2) a man riding a wave on a surfboard in the ocean. (p=0.004277)



- 0) a man riding a skateboard up the side of a ramp . (p=0.008467)
- 1) a man riding a skateboard on a ramp . (p=0.001182)
- 2) a man riding a skateboard up the side of a cement ramp . (p=0.000948)



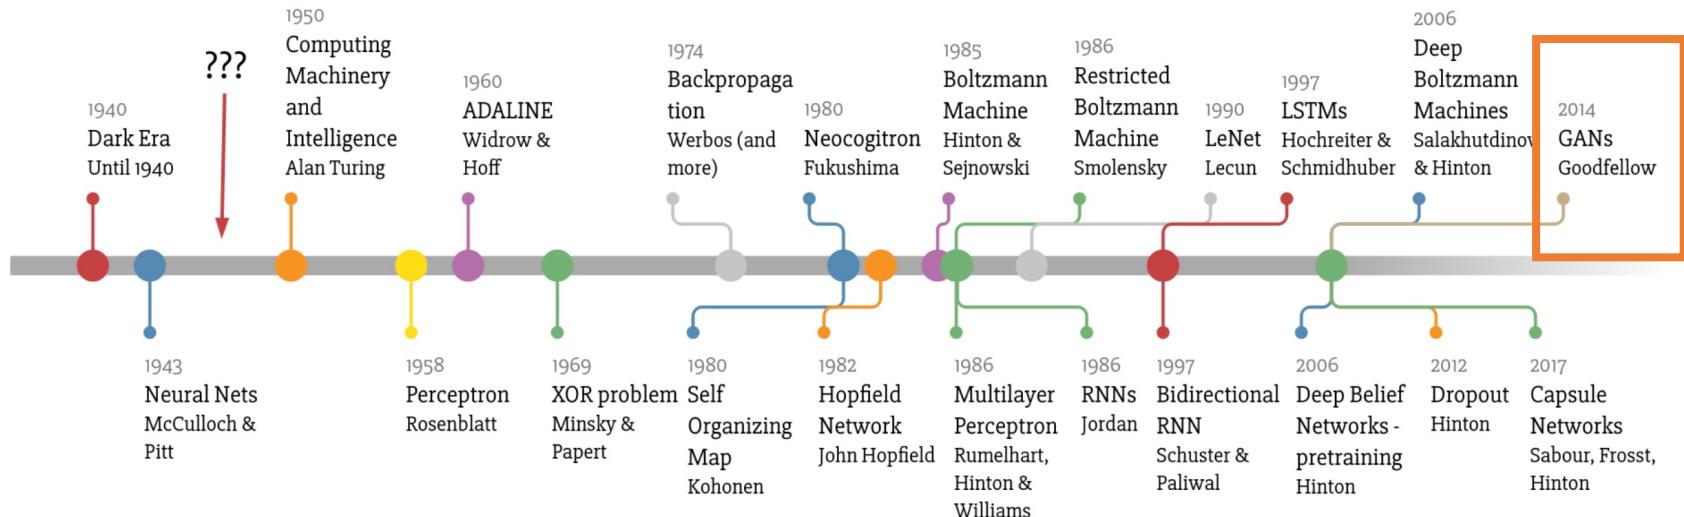
- 0) a ct showing a acute subdural hemorrhage with midline shift. (p=0.001608)
- 1) a ct showing a acute hemorrhage with left to right midline shift. (p=0.004277)
- 2) a ct showing a subdural hemorrhage with ventricular effacement. (p=0.021727)



- 0) a frontotemporal infarction with mild ventricular effacement. (p=0.001608)
- 1) a left frontotemporal infarction with mild left ventricular effacement. (p=0.004112)
- 2) a left fronto temporal infarction with ventricular effacement. (p=0.012928)

Feng, Rui, et al. "Deep learning guided stroke management: a review of clinical applications." *Journal of neurointerventional surgery* 10.4 (2018): 358-362.

# Deep Learning Timeline

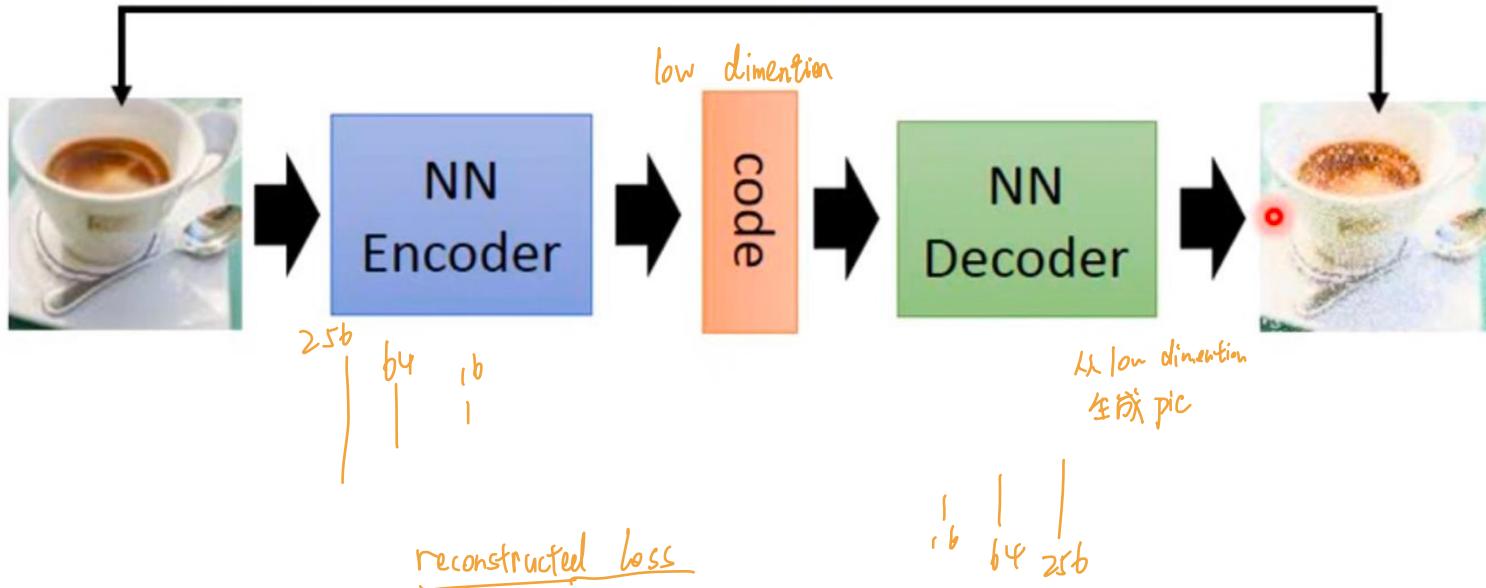


# Review: Auto-encoder

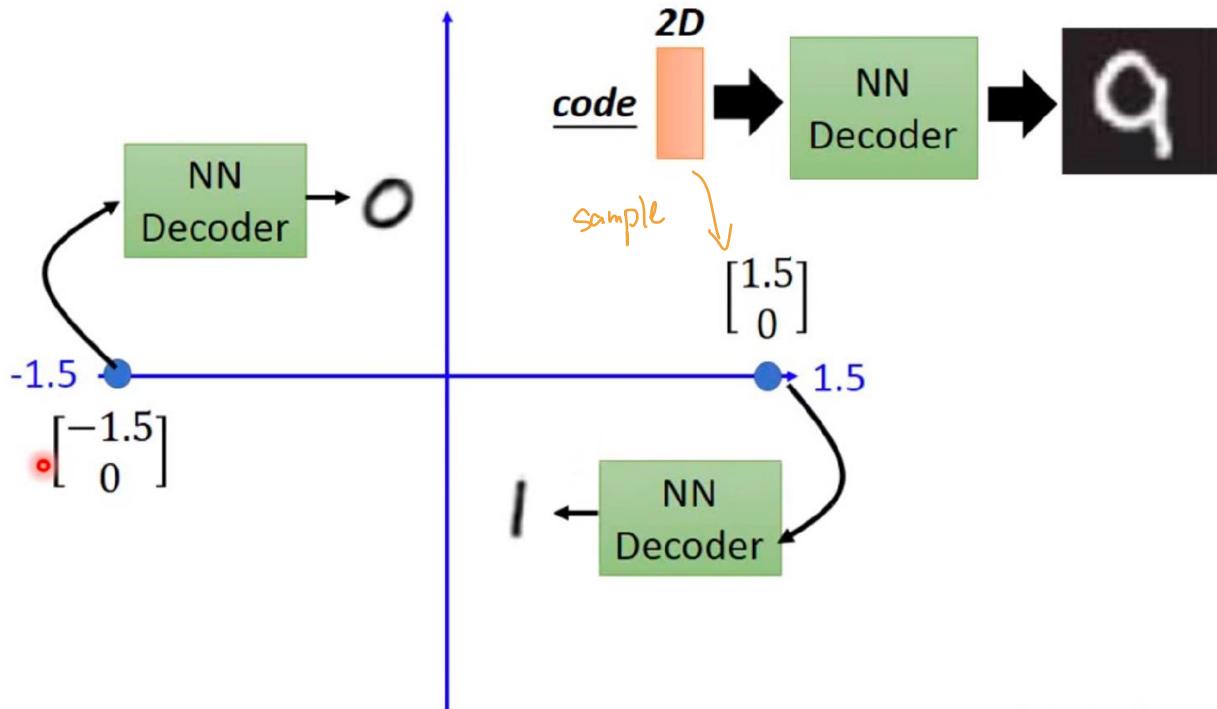
↪ unsupervised learning



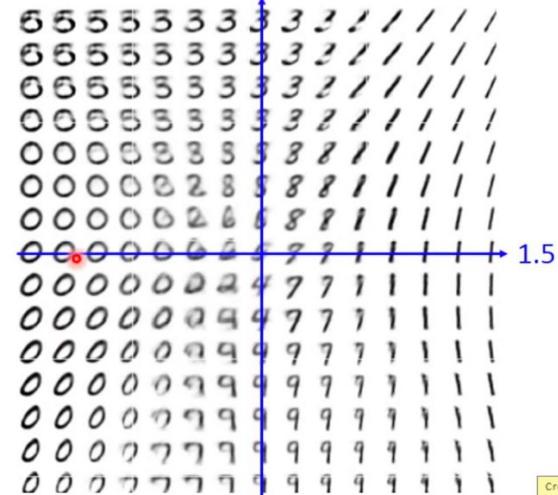
As close as possible



# Review: Auto-encoder



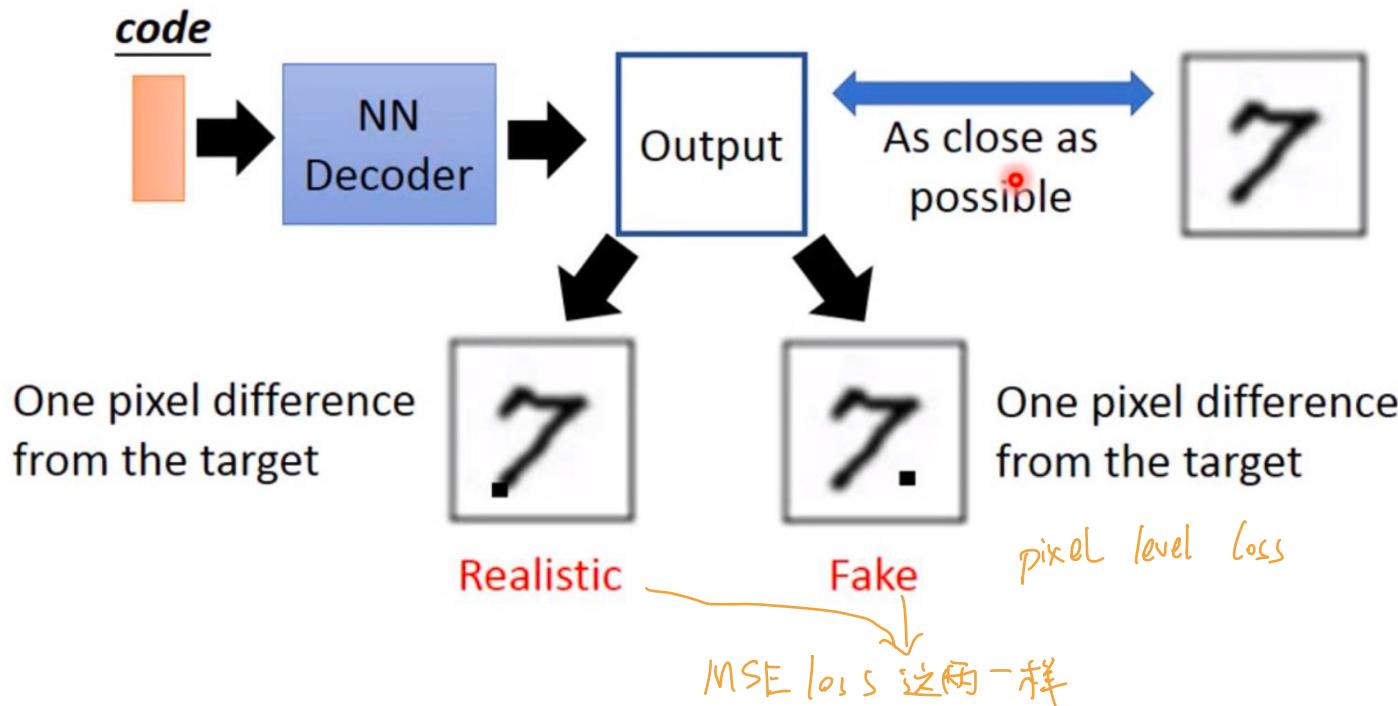
每个数字在 latent 中表达了



# Problems of VAE

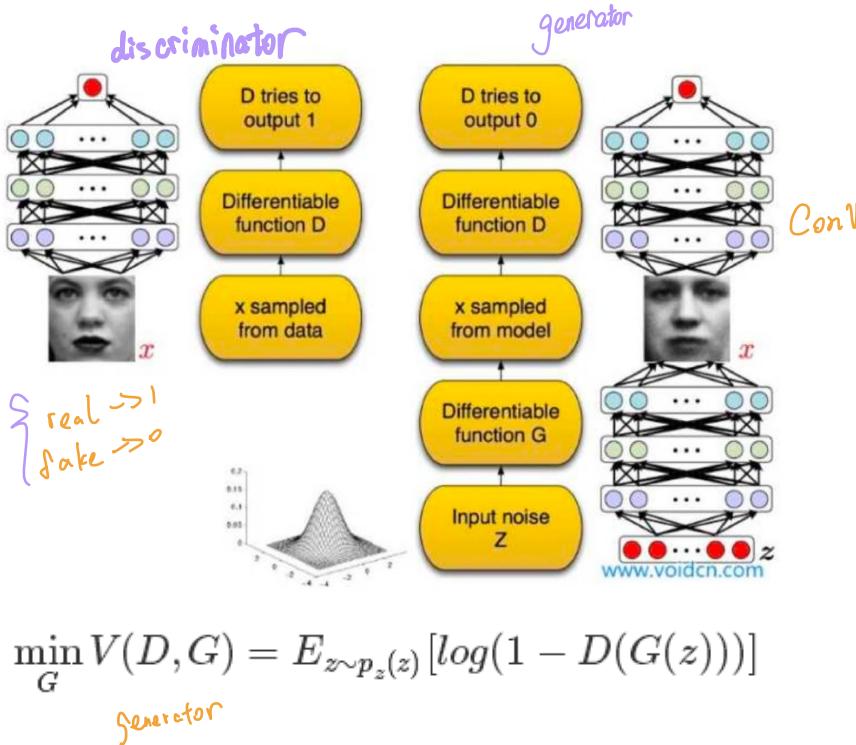
variational

- It does not really try to simulate real images



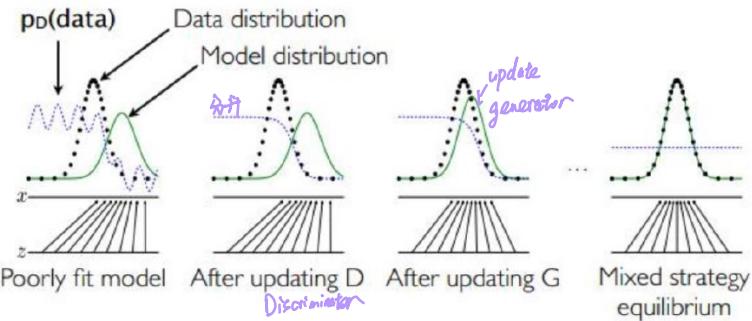
# Generative Adversarial Nets

$$\min_G \max_D V(D, G) = \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log(1 - D(G(z)))].$$



$$\min_G V(D, G) = E_{z \sim p_z(z)} [\log(1 - D(G(z)))]$$

*generator*



$$p_g(x) = p_{\text{data}}(x)$$

**Algorithm 1** Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator,  $k$ , is a hyperparameter. We used  $k = 1$ , the least expensive option, in our experiments.

for number of training iterations do

for  $k$  steps do

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Sample minibatch of  $m$  examples  $\{x^{(1)}, \dots, x^{(m)}\}$  from data generating distribution  $p_{\text{data}}(x)$ .
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[ \log D(x^{(i)}) + \log (1 - D(G(z^{(i)))) \right].$$

end for

- Sample minibatch of  $m$  noise samples  $\{z^{(1)}, \dots, z^{(m)}\}$  from noise prior  $p_g(z)$ .
- Update the generator by descending its stochastic gradient:

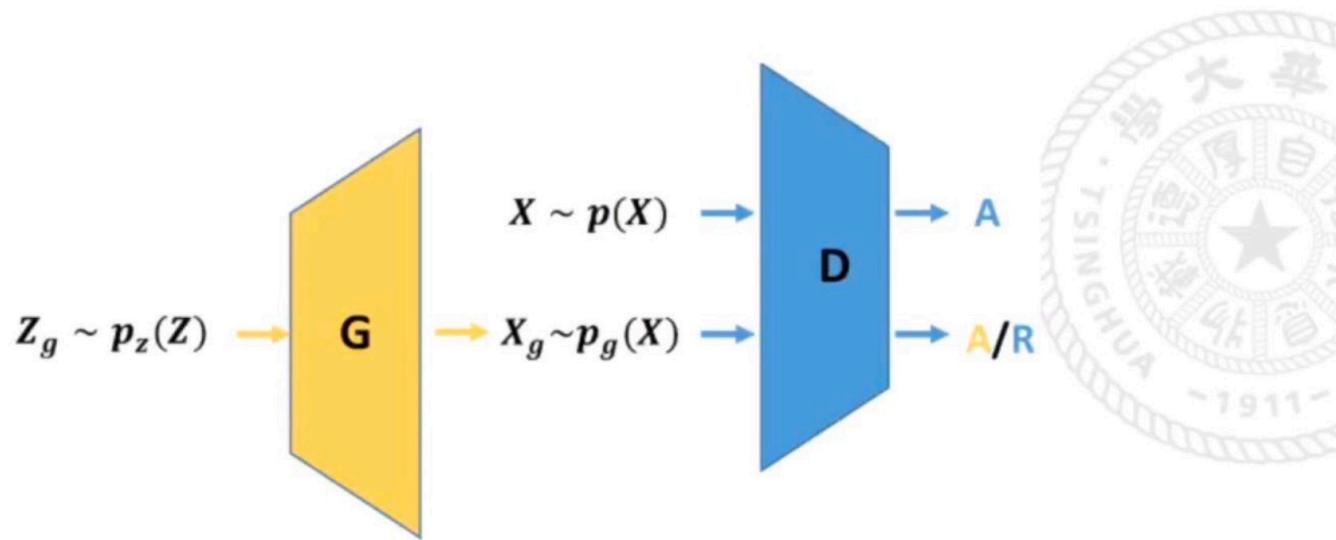
$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

# GANs (Goodfellow et al. [2014])

- Given  $p(x)$ , learn  $p_g(x) \approx p(x)$
- A two-player game



The global optima is  $p_g(x) = p(x)$ .

# The Two-player Formulation is Restricted

- $G$  and  $C$  may not be optimal at the same time. Here is an example of good  $C$  with poor  $G$  (Salimans et al. [2016]).

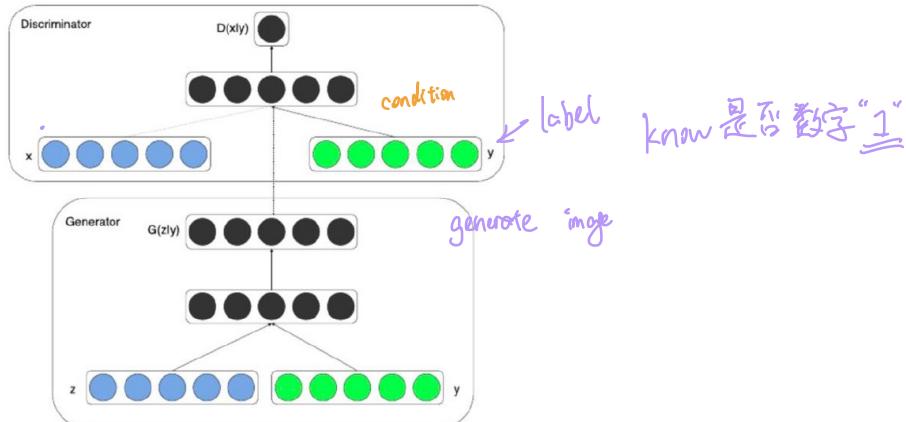


- $G$  cannot control the semantics of the generated samples.

# Generative Adversarial Nets

$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x})] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z})))].$$

# Conditional Generative Adversarial Nets



$$\min_G \max_D V(D, G) = \mathbb{E}_{\mathbf{x} \sim p_{\text{data}}(\mathbf{x})} [\log D(\mathbf{x}|y)] + \mathbb{E}_{\mathbf{z} \sim p_z(\mathbf{z})} [\log(1 - D(G(\mathbf{z}|y)))].$$

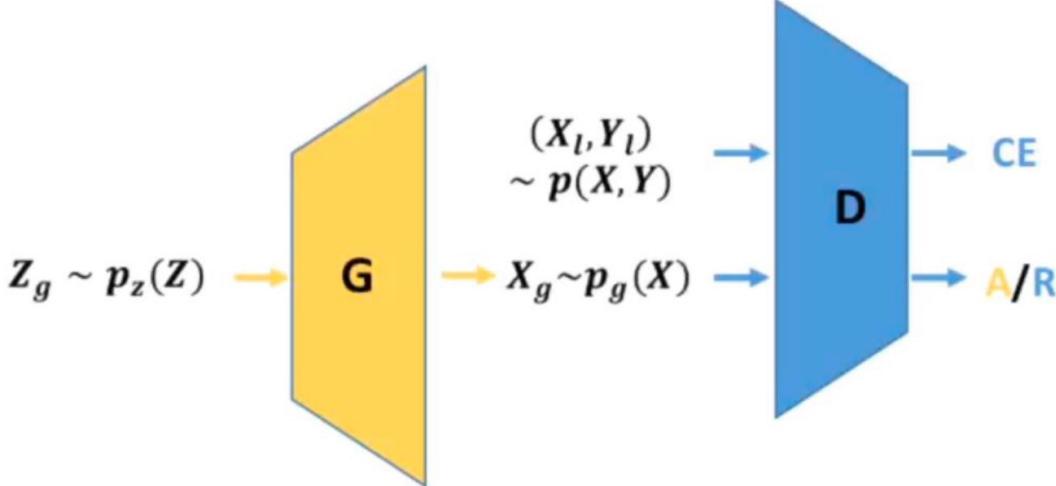


Figure 2: Generated MNIST digits, each row conditioned on one label

User tags + annotations	Generated tags
montanha, trem, inverno, frio, people, male, plant life, tree, structures, transport, car	taxi, passenger, line, transportation, railway station, passengers, railways, signals, rail, rails
food, raspberry, delicious, homemade	chicken, fattening, cooked, peanut, cream, cookie, house made, bread, biscuit, bakes
creek, lake, along, near, river, rocky, treeline, valley, woods, waters	water, river
people, portrait, female, baby, indoor	love, people, posing, girl, young, strangers, pretty, women, happy, life

Table 2: Samples of generated tags

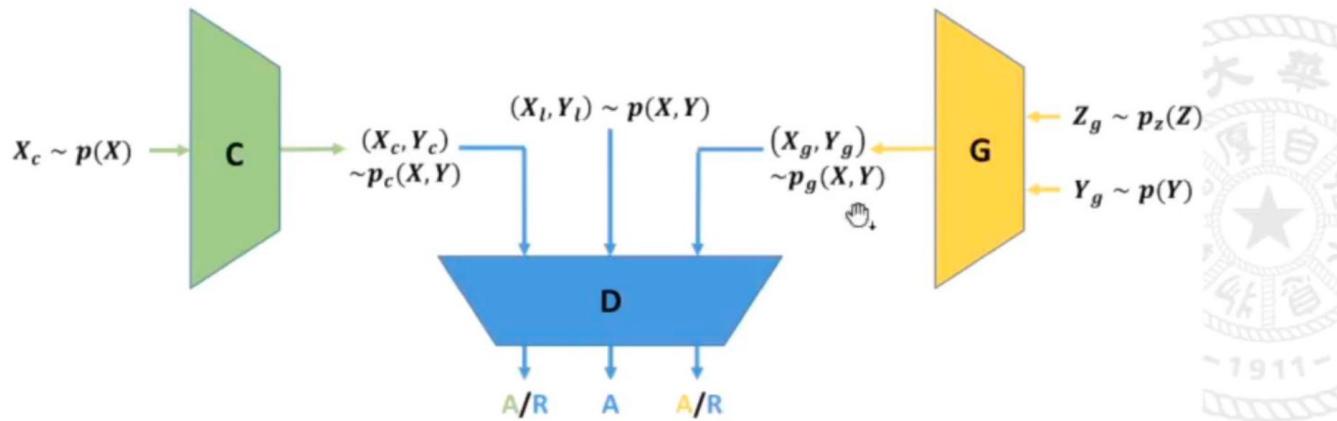
- Given partially labeled data, learn  $p_c(y|x)$  and  $p_g(x)$
- A two-player game with parameter-sharing



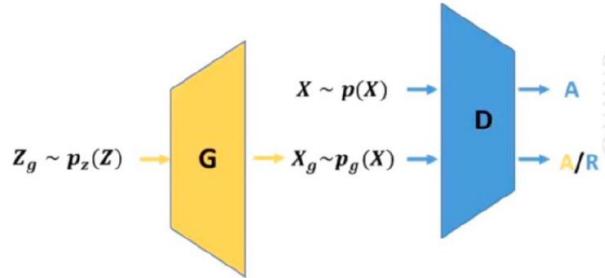
where  $D(x)$  outputs units of dimension  $|\mathcal{Y}|$  or  $|\mathcal{Y}| + 1$ .

# Three-player Game

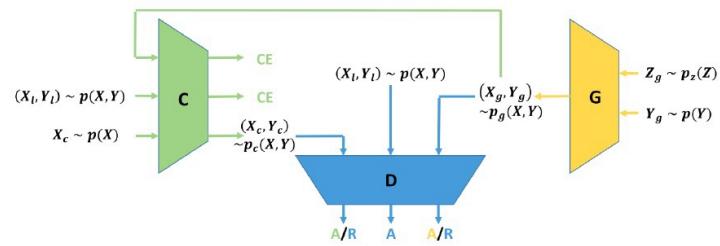
- Given partially labeled data, learn  $p_c(y|x)$  and  $p_g(x|y)$



where  $D(x)$  outputs a binary unit to justify  $(x, y)$  pairs. We call it Triple-GAN (three players and three distributions of interests).



$$\min_G \max_D U(D, G) = E_{x \sim p(x)}[\log(D(x))] + E_{z \sim p_z(z)}[\log(1 - D(G(z)))] ,$$



$$\begin{aligned} \min_{C,G} \max_D U(C, G, D) = & E_{(x,y) \sim p(x,y)}[\log D(x,y)] + \alpha E_{(x,y) \sim p_c(x,y)}[\log(1 - D(x,y))] \\ & + (1 - \alpha) E_{(x,y) \sim p_g(x,y)}[\log(1 - D(G(y,z),y))] , \end{aligned}$$

Add supervised loss (cross-entropy loss) to C

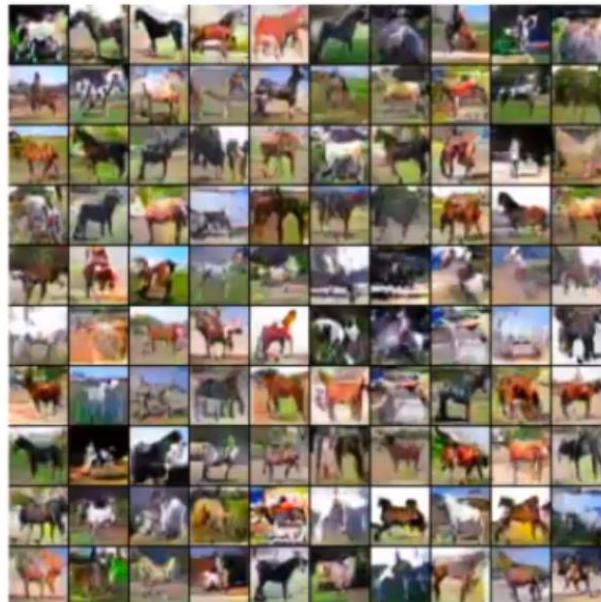
$$\mathcal{R}_{\mathcal{L}} = E_{(x,y) \sim p(x,y)}[-\log p_c(y|x)]$$

$$\begin{aligned} \min_{C,G} \max_D \tilde{U}(C, G, D) = & E_{(x,y) \sim p(x,y)}[\log D(x,y)] + \alpha E_{(x,y) \sim p_c(x,y)}[\log(1 - D(x,y))] \\ & + (1 - \alpha) E_{(x,y) \sim p_g(x,y)}[\log(1 - D(G(y,z),y))] + \mathcal{R}_{\mathcal{L}} . \end{aligned}$$

# Good $C$ Results in Good $G$



(a) Automobile



(b) Horse





Figure: Same  $y$  for each row. Same  $z$  for each column.

Style

# Home work: python + numpy

<https://cs231n.github.io/python-numpy-tutorial/>